


# Vector Databases

# Data Ingestion Pipeline: Introduction


# Vector Database

## Introduction



# Vector Database

## Data Types



# Vector Database

## Additional Resources

The screenshot displays the LangChain website interface. On the left, a sidebar navigation menu includes sections like Introduction, Tutorials, How-to guides, Conceptual guide, Ecosystem (LangSmith, LangGraph, LangServe), Versions (Overview, Release Policy, Packages, v0.2), and Security. The main content area features several large cards representing different components:

- LangSmith**: Observability.
- LangServe**: Deployments, Chains as Rest APIs (Python).
- Templates**: Reference Applications (Python).
- LangChain**: Cognitive Architectures, Chains, Agents, Retrieval Strategies (Python, JavaScript).
- LangChain-Community**: Integrations Components, Model I/O (Model, Prompt, Example Selector, Output Parser), Retrieval (Retriever, Document Loader, Vector Store, Text Splitter, Embedding Model), Agent Tooling (Tool, Toolkit) (Python, JavaScript).
- LangChain-Core**: Protocol, LCEL - LangChain Expression Language (Parallelization, Fallbacks, Tracing, Batching, Streaming, Async, Composition) (Python, JavaScript).


On the right, there are additional links for Tutorials, How-to guides, Conceptual guide, API reference, Ecosystem (LangSmith, LangGraph, LangServe), Additional resources (Security, Integrations, Contributing), Debugging, Playground, Evaluation, Annotation, and Monitoring.

Source: <https://python.langchain.com/>

# Data Ingestion Pipeline: Data Source and -Loading

# Vector Database


Data Source



# Vector Database


## Data Loading

- Hundreds of different data sources are supported by LangChain
- DataLoader returns list of LangChain documents
- Documents have two attributes
  - Metadata
  - page\_content



# Vector Database


## Data Loading



# Data Ingestion Pipeline: Data Chunking

# Vector Database

## Data Chunking




# Vector Database


## Data Chunking

### What is Data Chunking?


- Dividing larger pieces of information into smaller, manageable units
- These units called „chunks“
- Required to fit model context window
- Chunks should be:
  - Small
  - Semantically meaningful



Split



Combine




Overlap



# Vector Database

## Data Chunking: Chunking Approaches

### Overlap




Chunk 1    Chunk 2



**Fixed Chunk-Sizes**  
**Identical**  
**pre-defined**

**Structure-Based Chunk-Sizes**

- e.g. chat messages should be consistent, no mix of users and chunks



- Sentence 1 and 2 are very similar → same chunk
- Sentence 3 different → new chunk

**Semantic Chunking**


- based on semantic similarity
- e.g. when semantic break is observed

# Vector Database


Data Chunking: Splitter Types




Text



JSON



HTML



Code


# Vector Database

## Data Chunking: Splitter Types

- `chunk_size`...defines maximum size of chunks [characters]
- `chunk_overlap`...possible overlap of max 5 characters

The quick brown fox  
jumps over the lazy  
dog.\n This is a simple  
example to show text  
splitting.\n.

```
RecursiveCharacterTextSplitter(  
    chunk_size=20,  
    chunk_overlap=5  
    separators=[“\n”, “ “, “”]  
)
```



# Vector Database

## Data Chunking: Best Practices for Chunk Sizes

Texttyp	Granularität	Empfohlene Chunk- Größe (Tokens)	Empfohlener Overlap (Tokens)	Begründung
<b>FAQs / Kurze Q&amp;A</b>	Hoch (Fein)	<b>128 - 256</b> Tokens	<b>10 - 30</b> Tokens	Ermöglicht hochpräzise Antworten; minimiert "Rauschen" (irrelevante Infos).
<b>Code-Snippets / Logs</b>	Hoch (Fein)	<b>200 - 400</b> Tokens	<b>50 - 80</b> Tokens	Kurz genug, um eine einzelne Funktion/Logik-Einheit abzubilden; Overlap sichert Kontext (z.B. Funktionssignatur).

# Vector Database

## Data Chunking: Best Practices for Chunk Sizes

Technische Dokumente (z.B. Specs, Ingenieursberechnungen)	Mittel	500 - 800 Tokens	100 - 150 Tokens	Genug Platz für Rechenschritte, Formeln oder einen kompletten technischen Absatz; Overlap bewahrt den Fluss.
Allgemeine Artikel / Webseiten (Standard)	Mittel	512 - 1024 Tokens	100 - 200 Tokens	Dient als <b>ausgewogener Startpunkt</b> für die meisten Anwendungsfälle; sichert einen umfassenden Kontext.
Forschungspapiere / Verträge (Langform, hohe Kohärenz)	Niedrig (Grob)	1024 - 2000 Tokens	200 - 400 Tokens	Erlaubt die Aufnahme längerer Argumentketten, gesamter Abschnitte oder juristischer Klauseln, um den <i>globalen</i> Kontext zu erhalten.

# Vector Database

Data Chunking: Tokenization

The quick brown fox jumps over the lazy dog.



The

quick

brown

fox

jumps

over

the

lazy

dog

.



Tokenization

791

4062

14198

39935

35308

927

279

16053


5679

13

# Vector Database

## Data Chunking: Context Window


- Embedding model works with tokens, NOT words
- Model can cover only specific sequence lengths
- Too long text (longer than context window) will be truncated



# Data Ingestion Pipeline: Embeddings

# Vector Database


Embeddings: Introduction



# Vector Database

Embeddings: What?

- Conversion of text data into numeric vectors
- Each word / sentence is represented as vectors
- Vector has „low“ number of dimensions



# Vector Database

Embeddings: What?

Words    “Embedding” → Vectors

All  
data  
in  
deep  
learning  
must  
be represented  
represented  
as

We often call this embedding a word, which invites you to think of these vectors


▶ ▶ ⏪ 13:39 / 27:13 • Word embeddings > 🔍 ⚙ ☰

Source: <https://www.youtube.com/watch?v=wjZofJX0v4M&t=814s>

# Vector Database

Word Embeddings: What is it?

- Convert words to numbers
- Representation of words as unique tensors in high-dimensional space
- Relationships to other words are captured
- Ideally similar words are close
- Usually Deep Learning applied to get embeddings
- Embeddings represent meaning




Word Embeddings represent words  
as low-dimensional vectors  
in mathematical space and  
capture their semantic and  
syntactic meaning.

# Vector Database

Embeddings: Why?

- Semantic representation
  - capture meaning of data
  - enable comparison and analysis
- Lower dimensionality
  - computational complexity is reduced
  - high-dimensional data can be represented in lower dimensions
- Reusability
  - usable across different applications



Source: <https://nlp.stanford.edu/projects/glove/>


# Vector Database

From Words to Tensors

Input sentence



Tokenization



Tensor

# Vector Database

## Word Embedding Approaches

One-Hot Encoding

Frequency-Based

Neural Network

# Vector Database

One-Hot Encoding

Index:	0	1	2	3	4	5
Word:	It	was	a	bright	cold	day

	0	1	2	3	4	5
It	1	0	0	0	0	0
was	0	1	0	0	0	0
a	0	0	1	0	0	0
bright	0	0	0	1	0	0
cold	0	0	0	0	1	0
day	0	0	0	0	0	1

# Vector Database


## One-Hot Encoding - Problems

### Problems

- Curse of dimensionality → memory issues
- Matrix very sparse
- Words are isolated from each other
- All words have the same distance to each other

# Vector Database


## Word Embedding Approaches



- Very similar to OHE
- Gets count of words in document
- Term-Frequency/Inverse Term Freq.
- Gets count of words in document AND corpus
- Words frequent in a doc → important
- Words frequent in corpus → not important
- Gets similarity of words

# Vector Database


## Word Embedding Approaches



# Vector Database


## Neural Network based Embeddings

- Aim to
  - Capture context / meaning
  - Capture similarity to other words
  - Reduce dimension
  - Avoid memory issues
- Developed based on Neural Networks



# Vector Database

Word2Vec: Continuous Bag of Words




Independent Features	Dependent Feature“
["It", "was", "bright", "cold"]	"a"
["was", "a", "cold", "day"]	„bright“
...	

# Vector Database

Word2Vec: Continuous Bag of Words Model

Inputs




# Vector Database

Word2Vec: Skip Gram

Inputs


Model

Outputs



# Vector Database


## Word Embedding Approaches



# Vector Database

## GloVe


- Global Vectors for Word Representations
- Jeffrey Pennington, Richard Socher, and Christopher D. Manning. 2014. [GloVe: Global Vectors for Word Representation](#)
- based on co-occurrence matrix of words in a corpus, which counts how often words appear together in the same context.
- constructs a matrix of word co-occurrence counts and then factorizes this matrix to obtain word embeddings
- factorization based on singular value decomposition (SVD)
- resulting embeddings are dense, low-dimensional vectors
- Encode words as vector of other words



Source: <https://nlp.stanford.edu/projects/glove/>

# Vector Database

## Word Embedding Approaches




# Vector Database

## BERT

- Bidirectional Encoder Representations from Transformers
- Developed by Google in 2018
- Pre-trained word embedding
- Based on Transformers
- Applies „masked language modeling“ – masking some words in sentence and learn to predict them
- Applies „next sentence prediction“ – model predicts whether two sentences are similar in a text
- Original variants: BERT-base (110m parameters, 440MB) and BERT-large (340m parameters, 1.3GB)
- Other variants: RoBERTa, ALBERT, ELECTRA, ...

# Vector Database

## Word Embedding Approaches



# Vector Database

GPT

- Generative Pre-trained Transformers
- Developed by OpenAI
- Not strictly a word embedding, but contextualized word embedding
- Unique embedding for each occurrence of a word based on surrounding words in text
- Applies Transformer architecture
- GPT-3 has 175 billion parameters



# Vector Database

Difference Embedding Model vs. Large-Language Model

Parameter	Embedding Model	LLM
Base architecture	transformers	transformers
Process	texts, words, ... → numerical vectors	predict next words
Target	find semantic similarities of texts	generate outputs depending on context, e.g. next words
Applications	semantic search, clustering, representations for ML	text generation, QA systems, chatbots, translations, code generation


# Vector Database

Difference Embedding Model vs. Large-Language Model

Parameter	Embedding Model	LLM, LMM
Inputs	Text, words, sentences, images, ...	Text, words, sentences, images, ...
Output	vector	human-readable text/code
Focus	representation of data	processing and generation of data
Model Size	smaller, more specific (narrow AI, e.g. sentence transformers)	larger, e.g. GPT, Llama, ...
based on	pre-trained language models, uses architecture only for vector creation	uses transformers

# Vector Database

Difference Embedding Model vs. Large-Language Model




Source: [https://www.researchgate.net/figure/An-illustration-of-the-BERT-model-The-model-is-predicting-the-masked-word-brown\\_fig5\\_347822270](https://www.researchgate.net/figure/An-illustration-of-the-BERT-model-The-model-is-predicting-the-masked-word-brown_fig5_347822270)

# Vector Database

Embeddings: How?

Word Embeddings



Sentence Embeddings



# Vector Database

Embeddings: Which types are available?

Type	Model	Provider	Price	Vector Size
Online	text-embedding-3-small	OpenAI	0.02\$ / 1M tokens	1536
Online	text-embedding-3-large	OpenAI	0.13\$ / 1M tokens	3072
Online	mistral-embed	MistralAI	0.10\$ /1M tokens	1024
Offline	all-MiniLM-L6-v2	Open Source	---	384
Benchmark: <a href="https://huggingface.co/spaces/mteb/leaderboard">https://huggingface.co/spaces/mteb/leaderboard</a> ...				

# Vector Database

Embeddings: Factors to consider



Price



Speed



Off-/Online



Benchmark  
Performance


# Vector Database

Coding: Embedding GloVe closest words

- Find closest words

```
get_closest_words_from_word('chess')  
✓ 7.5s  
[('chess', 0.0),  
 ('backgammon', 4.379469394683838),  
 ('grandmasters', 4.56368350982666),  
 ('grandmaster', 4.613785743713379),  
 ('scrabble', 4.677640438079834)]
```

- Find word analogies




```
analogy = get_word_analogy(word1='sister',  
                             word2='brother',  
                             word3='nephew')  
analogy  
✓ 7.3s  
'niece'
```

# Vector Database

Coding: Word Cluster


- Given some categories
- Find words for the categories
- Check if they are „close“ (similar)



# Data Ingestion Pipeline: Data Storing

# Vector Database

## Data Storing



# Vector Database

Data Storing: What is a vector database?



A vector database stores [low-dimensional data](#) (embeddings) for fast querying and similarity analysis.

## Features

- Special type of database
- Allows to store, manage, and query data which is represented in geometric formats
- Enables similarity search, clustering, real-time analytics

# Vector Database

Data Storing: Why is a vector database needed?



Typical  
Databases



SQL

Structured Data

# Vector Database

Data Storing: Text Querying

Task: Which text is most similar?



Input Text Prompt: „Please provide a book on geometry“



Music



Sports



Math



History



Arts

Vector  
DB with  
Embeddings



Math



# Vector Database

Data Storing: Image Querying

Task: Which image is most similar to a text prompt?



Prompt: Which picture shows a brown wolf?



Most similar document

# Vector Database


Data Storing: Vector DB Providers



# Data Ingestion Pipeline: Data Querying

# Vector Database

Data Querying: Text Querying




Practical Implementation

```
collection.query(query_texts=[“This is my input text”])
```

# Vector Database

Data Querying: Image Querying




```
query_list = ["../data/dogs/akita_1.jpg"]
query_result = chroma_collection.query(
    query_images = query_list,
    n_results=3,
    include=['documents',
             'distances',
             'metadatas', 'data',
             'uris'],)
```

Result 1: ../data/dogs/akita\_3.jpg  
with distance: 0.17

# Vector Database

## Data Querying: Image Querying 2



```
query_list = ["dog in grassland"]
query_result = chroma_collection.query(
    query_texts = query_list,
    n_results=3,
    include=['documents',
        'distances',
        'metadatas', 'data',
        'uris'],)
```

Query: dog in grassland Result 0:  
./data/dogs/mastiff\_1.jpg  
with distance: 0.85

# Data Ingestion Pipeline: Similarity Search


# Vector Database

## Similarity Search

- Vector DB needs to analyze similarity of query-embedding compared to document embeddings.
- Approaches:
  - Cosine Similarity
  - Maximum Margin Relevance

# Vector Database

## Similarity Search



$$dist = \sqrt{(x_1 - y_1)^2 + (x_n - y_n)^2}$$

For an embedding vector of 768 embeddings, there are 768 distance terms


Example: word embeddings reduced to 2 dimensions

# Vector Database

Similarity Search

Imagename	Embedding					
dog1	0.3	0.02	0.8	0.6	...	0.4
dog2	0.1	0.52	0.7	0.6	...	0.4
	...					
dogN	0.3	0.62	0.9	0.2	...	0.3


Vector Database



# Vector Database

## Similarity Search: Cosine Similarity


- Measures similarity between Embedding-Vectors based on angle  $\theta$ .
  - Vectors maximally dissimilar  
→ vectors perpendicular ( $\theta = 90^\circ$ )
  - Vectors completely similar  
→ vectors parallel ( $\theta = 0^\circ$ )



# Vector Database

Similarity Search: Cosine Similarity

- Only the angle defines the similarity
- NOT the euclidean distance or magnitude of a vector
- Example
  - A: "The cat sleeps."
  - B: "The feline slumbers peacefully on the soft cushion."
  - C: "Trees grow leaves in spring."
  - D: "Fish swim in the ocean."



# Vector Database


Similarity Search: Maximum Margin Relevance

- Approach: reduce redundancy while maintaining relevance and diversity
- Redundancy...similar vectors
- Relevance...how closely do query and documents match
- Avoid clustering effect



Topic: Renewable Energies

What are the main types of renewable energy sources and how do they work?



# Data Ingestion Pipeline: Retrieval-Augmented Generation

# Data Ingestion Pipeline

Retrieval-Augmented Generation

