Question 1:

You are developing a compliance monitoring application that flags internal company reports containing employee names using Azure AI Language's Personally Identifiable Information (PII) detection feature.

Which PII category should you configure to detect employee names?

- ○ **Phone Number**
- ○ **DateTime**
- ◉ **Person**
- ○ **Age**

Correct

Mark for review

You are configuring the Azure AI Language Question Answering service to build a knowledge base by importing FAQ documents.

You need to identify the types of data that can be extracted during the import process.

Which types of data will be extracted?

Your answer is correct

✓ **Formatted text, URLs, and bulleted and numbered lists only**

Explanation

The service processes structured text content commonly found in FAQs, making it suitable for creating relevant question-answer pairs.

Question 3:

Scenario:
You are tasked with deploying an Azure AI Language solution to an environment that does not have internet connectivity. The solution should continue to operate offline while ensuring that the model and related components are properly deployed.

Which approach should you take for deploying the solution in this scenario?

- ○ **Download the model and deploy it to an offline virtual machine**
- ◉ **Deploy the solution to a Docker container hosted on a local machine**
- ○ **Use an Azure AI Services standard instance in a region with no internet connectivity**

Correct

Question 4:

You are building an app that will analyze meeting recordings and identify who is speaking at each moment in time.

You need to configure a voice profile for the app.

Which type of voice profile should you use?

- ○ **Text-Dependent Verification**
- ○ **Text-Independent Verification**
- ◉ **Speaker Identification**

Correct

Question 5:

You are building a mobile app that allows users to scan street signs and read the text aloud.

You need to recommend an Azure service that can recognize text from images with minimal development effort.

Which service should you recommend?

- ○ **Azure AI Custom Vision**
- ○ **Azure AI Document Intelligence**
- ◉ **Azure AI Vision**
- ○ **Azure AI Face Service**

Correct

Question 6:

You are developing a mobile app that allows users to scan and upload receipts. The app needs to extract specific information such as purchase dates, total amounts, and merchant names from the scanned receipts automatically.

Which Azure service should you use?

- ○ **Azure Application Insights**
- ◉ **Azure AI Document Intelligence**
- ○ **Azure AI Metrics Advisor**
- ○ **Azure AI Language**

Correct

You are working on a question-answering solution and configuring the API call for adding synonyms. However, you encounter an error when executing the following call:

```
"synonyms": [
    {
        "alterations": [
            "fix problems",
            "troubleshoot",
            "#diagnostic"
        ]
    },
    ...
]
```

To ensure the API call is executed successfully, which action should you take?

- (●) **Remove any special characters from the call**
- ( ) **Modify the order of the synonyms**
- ( ) **Remove the synonyms entirely**
- ( ) **Remove any question and answer pairs from the call**

Question 9:

You are developing an app that will use Azure AI Services.

You need to identify the methods that can be used to authenticate to Azure AI Services.

Which two methods can you use? Each correct answer presents a complete solution.

- [x] **An Azure subscription key**
- [ ] **A service principal**
- [x] **Microsoft Entra ID**
- [ ] **A client certificate**

Question 10:

You are developing a content moderation feature for a social media application that allows users to upload photos.

You need to ensure that uploaded images are automatically checked for adult content, explicit material, and other unsafe visual elements. The solution must require minimal development effort and be scalable.

Which Azure AI service should you use?

- ( ) **Azure AI Face Service**
- ( ) **Azure AI Vision Spatial Analysis**
- (•) **Azure AI Vision Image Analysis**
- ( ) **Azure AI Custom Vision**

Question 11:

You are building an app that transcribes audio recordings from customer service calls using Azure AI's Speech-to-Text feature.

During testing, the transcriptions contain frequent substitution errors, especially when recognizing specific product names and customer names unique to your business.

You need to improve the transcription accuracy and reduce the Word Error Rate (WER).

What should you add to the training data?

- ○ **Overlapping speakers**
- ○ **People talking in the background**
- ◉ **Custom product and people names**

You are building a customer support analysis app using Azure AI Language's Sentiment Analysis feature. The app processes product reviews to determine the overall sentiment of the feedback.

You test the app with a document containing one strongly positive review sentence and several neutral sentences describing product specifications without emotional tone.

Which sentiment label will the app return for the document?

Your answer is incorrect

⊗ **Neutral**

Explanation

Although most sentences are neutral, Azure AI Language prioritizes emotional sentences when determining the overall sentiment. Since there is one strongly positive sentence, the document is not classified as "Neutral."

Correct answer

○ **Positive**

Explanation

Azure AI Language assigns the "Positive" label because even a single clearly positive sentence can influence the overall sentiment when all other sentences are neutral. The service interprets the document as having a positive sentiment due to the presence of positive feedback.

Question 13:

You are developing an orchestration workflow for Language Understanding (LUIS) in Azure. The solution must support multiple languages while minimizing the administrative overhead required for configuration.

Which approach should you use to handle different languages efficiently?

- (●) **Separate workflow projects**
- ( ) **Separate training jobs**
- ( ) **A new deployment**
- ( ) **Separate models**

Question 14:

You are building an application using the Azure AI Translator to translate content between languages. To enhance the capabilities of the application, you want to incorporate additional features of the Translator service.

Which three features are supported by the Azure AI Translator service? Select all that apply.

Your selection is incorrect

☒ **Intent Recognition**

Explanation

Intent recognition is part of the LUIS service and not available in the Translator service, which focuses primarily on translation and related linguistic capabilities.

Correct selection

☐ **Transliterate**

Explanation

Transliteration is a part of Azure Translator, particularly useful in scenarios where the script of the language needs to be changed without altering the meaning of the words.

Your selection is correct

☑ **Dictionary Lookup**

Explanation

This feature is available in the Azure AI Translator, allowing applications to fetch detailed word meanings and contextual information, enriching the user experience.

☐ **Entity Extraction**

Explanation

Azure Translator does not perform entity extraction. Instead, **Text Analytics** or other NLP services would be used to extract entities from the text.

You are building an app that will use Azure AI Custom Vision. The app will be deployed to a virtual machine in Azure.

You enable firewall rules for your Azure AI Services account.

You need to ensure that the app can access the service through a service endpoint.

What should you do?

○ **Grant access to an internet IP range**

Explanation

1. Allowing access from an **internet IP range** would expose the Azure AI service to the **public internet**, which **does not involve service endpoints**.
2. Granting IP access is useful for **specific internet-based services**, but it reduces security and **doesn't provide private connectivity** through a service endpoint.

Your answer is incorrect

⊗ **Assign a role-based access control (RBAC) role to the Azure AI Custom Vision resource**

Explanation

1. While **RBAC roles** manage access permissions for Azure resources, they do not configure network connectivity through **service endpoints**.
2. RBAC ensures **identity-based access control**, but enabling a service endpoint requires **network-level access configuration**, such as connecting a virtual network.

○ **Include an access token in the Authorization header**

Explanation

1. Including an **access token** in the Authorization header is a method of **application-level authentication** but **doesn't configure network-level access**.
2. This approach is unrelated to enabling service endpoints, which requires **virtual network configuration**, not token-based access.

Correct answer

○ **Grant access to a specific virtual network**

Explanation

1. Azure **service endpoints** extend your virtual network's private IP address range to Azure services, such as Azure AI Custom Vision.
2. **Firewall Rules** in Azure AI Services allow access only from approved virtual networks when configured properly.

Question 16:

You are developing an application that will extract the key concepts from a document using Azure AI Services.

Which endpoint should you use as part of the solution?

- ○ **Azure AI Vision API**
- ○ **Custom Named Entity Recognition (NER)**
- ◉ **Key Phrase Extraction**

You are creating a custom translation model using Azure AI's Translator service. You want to train the model with bilingual documents to ensure the model understands your domain-specific terminology and style.

Which approach should you take when preparing your training data?

Correct answer

○ **Be liberal**

Explanation

Being liberal means including diverse and extensive bilingual documents that reflect various contexts, terminology, and sentence structures. This broad coverage helps the translation model learn more effectively and perform better in real-world scenarios.

Your answer is incorrect

⊗ **Be strict**

Explanation

Being strict implies using only perfectly curated and narrowly defined examples. While quality data is important, excessive strictness could limit the amount of training data and reduce the model's versatility.

○ **Be restrictive**

Explanation

Being restrictive would limit the training data to only a narrow set of terms or contexts. This would constrain the model's learning and make it less adaptable to varied translation tasks.

Question 18:

You are developing an Azure-based application that processes customer service transcripts. The application must identify and redact sensitive information such as names, phone numbers, and email addresses.

Which categories should you configure in Azure AI Language PII detection to meet the application's requirements?

- ○ **Address, IP Address, and Credit Card Number**
- ○ **Organization, Email Domain, and Phone Code**
- ◉ **Name, Phone Number, and Email Address**
- ○ **Name, PersonType, and Location**

You are developing a multilingual conversational app using **Conversational Language Understanding (CLU)** from Azure AI Language services. The model you've created needs to support multiple languages, and you want to optimize its performance with minimal development effort.

What should you do to improve the performance of the model?

Your answer is incorrect

⊗ **Train the model with utterances in multiple languages and query it using only the model's primary language**

Explanation

This is inefficient because the model is designed to handle multiple languages. Limiting queries to only one language negates the advantage of using a **multilingual model**. The model's purpose is to understand and respond to queries in different languages, and restricting the query language to just one reduces the overall performance and flexibility of the app.

Correct answer

○ **Add more utterances for languages that are underperforming in the model**

Explanation

1. **Data-driven improvement**: The best way to improve the performance of a model, especially when it's multilingual, is to enhance its **training data**. By adding more utterances for languages where the model is underperforming, you ensure that the model learns a broader and more diverse set of language patterns, idioms, and possible user inputs.
2. **Minimizes development effort**: This approach doesn't require rearchitecting the model or adjusting the entire system. It is the simplest and most effective method for improving model performance.

- ◯ **Create separate models for each language**

Explanation

While creating separate projects for each language might improve performance in certain cases, it significantly increases **complexity** and **maintenance overhead**.

- ◯ **Configure the app to only query utterances in the language used to train the model**

Explanation

This solution undermines the multilingual capability of the app. The purpose of using a **multilingual model** is to support multiple languages simultaneously. Restricting the model to only process queries in a single language goes against the goal of having a global, diverse conversational app.

Question 20:

You plan to build an app that will transcribe large quantities of audio files using the Azure AI Speech service's batch transcription feature.

You need to recommend a storage solution for the audio files. The solution must minimize development effort.

What should you recommend?

- ○ **Azure Data Lake Storage**
- ◉ **Azure Storage**
- ○ **Azure Cosmos DB**
- ○ **Azure SQL Database**

You are configuring the Azure AI Language Question Answering service for a production-ready chatbot. To estimate the operational costs of the service, you need to identify which factors influence the service's pricing.

Which three parameters will directly impact the operational costs of the Azure AI Language Question Answering service?

Your selection is correct

☑ **The size and number of knowledge bases used**

Explanation

1. Azure charges based on the number and size of knowledge bases created. Larger knowledge bases require more storage and indexing, increasing costs.
2. **Example:** A global enterprise with multiple departments may maintain several large knowledge bases, leading to higher service fees.

Your selection is incorrect

☒ **The number of predefined answers in the knowledge base**

Explanation

Azure doesn't charge based on the number of predefined answers in the knowledge base. Costs depend on the size of the knowledge base in terms of data storage, not the number of answers.

Correct selection

☐ **The required service throughput**

Explanation

1. Service throughput refers to the processing capacity needed for the Question Answering service. Higher throughput ensures lower response times for concurrent requests, leading to additional costs.
2. **Example:** A chatbot used during peak shopping seasons may need a higher throughput tier to avoid delays, increasing costs.

☐ **The number of administrators managing the knowledge base**

Explanation

Azure pricing isn't affected by the number of administrators or editors managing the knowledge base. Admin management is part of the platform's built-in functionality and doesn't directly impact service costs.

Your selection is correct

☑ **The number of user queries processed per second**

Explanation

- Azure AI Language charges based on API usage, measured by the number of user queries processed. The more queries processed, the higher the cost. This aligns with the "pay-as-you-go" pricing model in Azure.
- **Example:** If a chatbot receives thousands of queries daily, this will significantly increase operational costs.

You are planning to build a chatbot using Azure AI Language's **Question Answering** service to help users answer FAQs.

You need to identify which scenarios are best suited for integration with the **Azure AI Language Question Answering** service.

Which three of the following scenarios should you select as suitable for use with this service? Each correct answer presents a valid solution.

- [ ] **When the chatbot is designed to answer questions based on real-time, dynamic data**
- [ ] **When the chatbot provides personalized answers based on real-time user inputs**
- [x] **When the chatbot must provide consistent, repeatable answers for specific questions**
- [x] **When you have a knowledge base that includes pre-defined answers to frequently asked questions**
- [x] **When you have a chatbot that only handles static, predefined information**

correct

You are designing a content management application that extracts important points from lengthy articles for quick reference using Azure AI Language.

Which Azure AI Language feature should you recommend to extract key phrases from text documents?

○ **Named Entity Recognition (NER)**

Explanation

- Named Entity Recognition (NER) detects specific entities like names, locations, dates, and organizations.
- While useful for identifying structured data, it does not extract general key phrases or summarize text.

Your answer is correct

✓ **Key Phrase Extraction**

Explanation

- Key Phrase Extraction identifies the most relevant phrases from text documents, highlighting important concepts and terms.
- Example: In the sentence "Azure AI Language provides advanced text processing capabilities," the key phrases would be "Azure AI Language" and "text processing capabilities."
- This feature is ideal for applications that need to extract key points, making it the correct answer.

○ **Language Summarization**

Explanation

1. Azure AI Language does not have a feature explicitly named Language Summarization.
2. The correct feature for summarizing documents is Abstractive Text Summarization, which creates a concise summary rather than extracting key

Question 24:

You are working on an Azure AI Translator custom model and aim to improve its translation quality. You need to achieve a **Bilingual Evaluation Understudy (BLEU)** score that demonstrates a high-quality translation.

What is the minimum BLEU score range required to indicate a good quality translation?

Correct answer

◯ **40 to 59**

Explanation

1. The BLEU score is used to measure how closely the machine-generated translations align with human reference translations. A score in this range indicates that the model has been trained well enough to produce translations that are close to the reference translations but may still contain some minor discrepancies.
2. A BLEU score in this range would be suitable for a variety of production-level scenarios, ensuring the translation system provides high-quality results.

◯ **80 to 100**

Explanation

Achieving such a high BLEU score generally requires a large volume of data and extensive fine-tuning, which may not always be necessary or practical for many applications, especially early in the development phase.

Your answer is incorrect

⊗ **60 to 79**

Explanation

Although achieving this range would show excellent performance, it is beyond the minimum threshold required for good-quality translations and not necessary to indicate the high-quality translation needed in many use cases.

○ **0 to 19**

Explanation

A BLEU score in this range signifies that the model has a lot of room for improvement and would not be considered for production-level tasks.

○ **20 to 39**

Explanation

Though this range might indicate moderate progress, the quality is not yet suitable for high-accuracy translation needs, especially in business or critical applications.

To evaluate how accurately the model identifies relevant intents or responses, you need a metric that calculates the ratio of correctly predicted positive results to the total predicted positive results.

Which evaluation metric should you use?

○ **BLEU (Bilingual Evaluation Understudy)**

Explanation

BLEU is used to evaluate the quality of machine-translated text by comparing generated translations to reference translations. It is unrelated to intent recognition or positive classification in CLU models.

Your answer is correct

● **Precision**

Explanation

- Precision measures the percentage of correctly identified positive instances out of all instances the model classified as positive. It is the most relevant metric when assessing how accurate the model's positive predictions are.

○ **F1 Score**

Explanation

The F1 score balances precision and recall, providing a harmonic mean. While useful for comprehensive evaluation, it is not specific to measuring only the correctly predicted positives relative to all predicted positives.

○ **Recall**

Explanation

Recall measures how well the model identifies all relevant positive instances out of the actual positive cases. It does not focus on the accuracy of predicted positives, making it unsuitable for the metric described in the question.

You are building a containerized application that uses Azure AI Services for text extraction from scanned documents. During development, the container fails to connect to the AI Services resource, returning a "Resource Mismatch" status message.

You need to resolve this issue and establish a successful connection.

What should you do?

Correct answer

○ **Confirm that the API key is for the correct resource type**

Explanation

1. The **"Resource Mismatch"** error occurs when the API key being used **does not match the type of Azure AI Services resource** (e.g., using a Computer Vision API key for Text Analytics).
2. Each AI service has its **own API key** and **endpoint**. Verifying that the correct **resource type** is used ensures proper authentication and service communication.

Your answer is incorrect

⊗ **Confirm that the API key is for the correct region**

Explanation

1. API keys are **region-specific**, meaning you must deploy resources in the same region where the API key was generated.
2. However, a **"Resource Mismatch"** error typically indicates an issue **unrelated to the region** and more related to the **resource type**.

○ **Upgrade the Azure AI Services resource to a higher tier**

Explanation

1. **Resource tier upgrades** improve performance, scalability, and quotas but **do not affect API key mismatches**.
2. The issue here is related to **authentication**, not **resource capacity** or **performance limitations**.

- ○ **Confirm that the Azure AI Services resource is online**

Explanation

1. If the resource were offline, the status message would indicate unavailability rather than a mismatch.
2. Resource status issues would result in connection timeout or service unavailability errors, not a mismatch message.

Question 2:

You are working with a custom language content model in Azure AI Video Indexer. During testing, you upload a text file containing the sentence: "Kubernetes is a new feature in Azure & the cloud." However, the sentence is discarded by the system.

You need to ensure that the model retains the sentence during processing.

What action should you take?

- ○ **Switch to a custom slate detection model**
- ○ **Remove the special characters from the entire text file**
- ○ **Retrain the model**
- ◉ **Replace the "&" character with "and"**

Question 3:

You are building a GPT-based chatbot to answer questions about your company. You plan to use the "Using your data" feature in Azure OpenAI to ground the model with relevant company data.

Which of the following file types can you use to ground the model? Select all that apply.

Your selection is correct

☑ **TXT**

Explanation

TXT files are ideal for grounding purposes because they contain raw text without any special formatting or structure, making it easier to process and use directly in the model's training.

Your selection is correct

☑ **HTML**

Your selection is incorrect

☒ **JSON**

Explanation

JSON is used for data exchange and may not be processed directly for grounding OpenAI models. The feature works better with text-based file formats that provide richer textual content.

Your selection is incorrect

☒ **XML**

Explanation

Although XML can store textual data, it is a structured format that may not be as straightforward to process as HTML, Markdown, or PDF files. It generally requires more preprocessing before it can be used effectively for grounding a model.

Your selection is correct

- [x] **Markdown (MD)**

Explanation

Markdown is often used for company documentation, making it an excellent source for grounding a model with company-specific data. The model can parse the Markdown content to understand and answer questions based on the company's materials.

Your selection is correct

- [x] **PDF**

Explanation

PDFs, especially those that contain textual data like company reports, manuals, and whitepapers, can be used to ground an AI model. When processed properly, the text content from PDFs can provide the AI with domain-specific knowledge.

You are building an app that leverages the Azure AI Video Indexer API to analyze meeting recordings from a company's internal collaboration tool. The app needs to search for images and mentions related to competitors during the meetings.

Which content model should you use to identify specific brands or competitors in the meeting recordings?

Your answer is incorrect

⊗ **Custom Visual Content**

Explanation

Although it's focused on visual content, it lacks the brand recognition capability necessary for detecting **brands or logos** in the way that the **Custom Brand** model does.

○ **Custom Language**

Explanation

While **Custom Language** can be helpful for transcription or keyword spotting, it doesn't have the specialized capabilities needed to detect and categorize **brands** or **visual images**. This would not effectively solve the problem of identifying competitors in meeting recordings.

Correct answer

○ **Custom Brand**

Explanation

Since the requirement is to detect **images** and **mentions** of competing companies, the **Custom Brand** model is the best fit as it can recognize branded content in both visual and audio forms.

○ **Custom Speech**

Explanation

You are building a solution using Azure AI Search, and you need to store normalized binary files as projections.

Which type of projection should you use to handle binary data efficiently?

○ **tables**

Explanation

"Tables" projection is used to handle data that is organized in tabular format, which is common for relational or structured datasets like those found in databases. While you can store structured data in tables, binary data such as files would not be effectively managed using the "tables" projection. This projection is not suitable for handling files or unstructured binary data.

Correct answer

○ **files**

Explanation

When working with binary files, the "files" projection type is the most suitable option. This allows Azure AI Search to store binary data in its native form, ensuring efficient retrieval and search operations for documents or media stored as binary objects. For example, when you're indexing files such as images, PDFs, or videos, using the "files" projection allows you to process and store them as binary objects, ensuring the data remains in its raw format for searching and retrieving.

Your answer is incorrect

⊗ **objects**

Explanation

- The "objects" projection is typically used for structured data where you expect to work with a series of key-value pairs or object-like structures. This projection would not be appropriate for binary data like files, as it's designed to store and query structured data in the form of objects.

1.

You are building an app that will utilize Azure AI Vision to analyze and detect animals in images.

Which type of model should you choose for this application?

- (●) **Object detection**
- ( ) **Image segmentation**
- ( ) **Image classification**
- ( ) **Object classification**

You are building a solution using **Azure AI Search** and want to execute the initial run of the indexer. Understanding the internal process of how Azure Cognitive Search ingests, processes, and indexes data is critical.

The correct answer involves knowing the sequence of steps Azure AI Search follows when processing data during indexing.

Correct answer

◯ **Document Cracking, Field Mapping, Skillset Execution, and Output Field Mapping**

Explanation

When an indexer in Azure AI Search runs for the first time, it performs these essential tasks:

1. Document Cracking:
   - This process extracts text and metadata from structured and unstructured data sources like PDFs, images, and JSON files.
   - For example, extracting text from a scanned invoice image.
2. Field Mapping:
   - Maps data from the source fields to the corresponding fields in the search index.
   - For example, mapping a "Product Name" field from a database to the

"ProductTitle" field in the Azure AI Search index.

3. Skillset Execution:
   - Executes AI-powered skills like language translation, sentiment analysis, or image recognition if a cognitive skillset is configured.
   - For example, applying OCR to scanned documents.
4. Output Field Mapping:
   - Stores processed data in the correct index fields after transformations and AI processing.
   - For example, saving recognized text from a document into the "ExtractedText" field in the index.

○ **Connecting to an Azure Data Source, Creating an Index Schema, and Running the Wizard to Create Objects and Load Data**

Explanation

1. These are initial configuration steps **before** running the indexer, not actions the indexer performs during execution.
2. The indexer assumes that the data source, index, and objects already exist before it runs.
3. Example: You create these settings in the Azure portal or with APIs prior to

Your answer is incorrect

⊗ **Creating a Data Source, Creating an Index, and Creating and Running the Indexer**

Explanation

1. These are **pre-configuration steps**, not part of the indexer's execution flow.
2. The indexer requires the data source and index to be configured **before it runs**, meaning this step happens outside the indexing process.

You are designing an application that analyzes images uploaded by users to detect dominant background colors using the Azure Image Analysis API.

Which of the following colors can the API return as a dominant background color?

○ **Coral**

Explanation

Though "Coral" is recognized in some CSS color schemes, it is not included in the list of supported dominant background colors returned by the Image Analysis API.

Correct answer

○ **Teal**

Explanation

"Teal" is one of the common web-safe colors supported by the Image Analysis API. The API evaluates image pixels and categorizes dominant colors using standard color definitions, including "Teal."

○ **Lavender**

Explanation

The Azure Image Analysis API uses a standardized list of common colors based on web-safe and CSS-defined color names. "Lavender" is not part of the predefined list for dominant background colors.

Your answer is incorrect

⊗ **Magenta**

Explanation

While "Magenta" is a well-known color, it is not included in the list of common background colors the API returns. The API sticks to commonly used web-safe colors.

Question 9:

You are deploying an Azure OpenAI service and plan to use your own data for model deployment. To ensure that the model can index your data sources effectively, which additional Azure service should you deploy?

- ○ **Language**
- ○ **Content Moderator**
- ◉ **Azure AI Search**
- ○ **Personalizer**

You are developing a customer support chatbot using Azure OpenAI and have enabled the "Using your data" feature to ground the model with company-specific content.

During testing, you find that some responses include irrelevant or loosely related documents, reducing the chatbot's answer accuracy.

Which parameter should you adjust to ensure only highly relevant documents are considered during responses?

Your answer is incorrect

⊗ **Retrieved Documents**

Explanation

Increasing the number of retrieved documents may result in more irrelevant results unless the filtering parameter is adjusted.

○ **File Name**

Explanation

The model does not consider file names when evaluating content relevance.

○ **File Type**

Explanation

Filtering based on file type affects file acceptance but does not influence document relevance in responses.

Correct answer

○ **Strictness**

Explanation

Setting the strictness parameter to a higher value ensures only documents closely related to the query are considered, improving response accuracy.

Question 11:

You are developing a web application that generates images based on user input. The application uses the DALL-E 3 model provided through the Azure OpenAI service.

Which HTTP body property must be included in the request to successfully generate images?

- ○ **The Azure OpenAI resource name**
- ◉ **The prompt**
- ○ **The deployment ID**
- ○ **The API version**

You are developing a document search solution using **Azure AI Search**. You want to allow users to filter search results by categories and display the total number of documents that match each category.

Which field attribute should you assign to the field?

Your answer is incorrect

⊗ **filterable**

Explanation

1. This attribute allows fields to be used in **filter expressions** to narrow down search results, such as applying a price filter.
2. However, **filterable fields** do not display hit counts by category, which is what **facetable** does.

○ **retrievable**

Explanation

1. Fields marked as **retrievable** are included in the search results and can be displayed in the output.
2. This attribute controls the visibility of fields but does **not support faceting or counting.**

○ **sortable**

Explanation

1. **Sortable fields** allow sorting of results based on field values (e.g., sorting by price or date).
2. Sorting does not provide hit counts or categorical groupings.

Correct answer

○ **facetable**

Explanation

1. **Facetable fields** allow Azure AI Search to group search results into categories and display hit counts for each category.
2. This enables users to see how many items belong to each category, making navigation and filtering easier.

Question 13:

You are deploying an Azure OpenAI Service resource.

You need to configure the network access settings to ensure that only applications within your Azure subscription can access the service securely.

Which network access configuration should you choose?

- ○ **Allow access from selected IP ranges and enable a network security group (NSG)**
- ○ **Allow access from all networks**
- ○ **Allow access from selected virtual networks and IP addresses**
- ◉ **Disable public network access and use a private endpoint**

Question 14:

You are developing an application that needs to identify the presence of people in a live video stream. Which Azure AI Vision feature should you utilize?

○ **Image Analysis**

Explanation

Image analysis provides insights like object detection, scene descriptions, and color analysis, but it doesn't specialize in tracking human movement or detecting the presence of people across a dynamic video feed.

Your answer is incorrect

⊗ **Face Detection**

Explanation

1. Face detection is designed to recognize and locate human faces within an image or video, but it doesn't necessarily provide a complete understanding of human presence or movement across an entire scene.
2. It is focused more on identifying the position of faces within images and videos rather than tracking general human presence in broader contexts (e.g., identifying multiple people or tracking their movements across a space).

Correct answer

○ **Spatial Analysis**

Explanation

Spatial Analysis is the correct Azure AI Vision feature for detecting the presence of people in a video feed. It provides capabilities for analyzing real-time video streams, including detecting and tracking human presence, movement, and other spatial elements in the scene. This feature is ideal for video-based surveillance or interactive applications where recognizing people and their actions is a requirement.

○ **Optical Character Recognition (OCR)**

Explanation

OCR is designed to extract text from images or video frames. It is not relevant to detecting people or analyzing human presence.

Mark for review

You are building an app that will use Azure AI Vision to analyze and classify images in order to build an image library of animals. You want to ensure that each image selected contains only one animal.

Which type of classification should you use to meet this requirement?

Your answer is incorrect

⊗ **Multiclass**

Explanation

Multiclass classification means that each image is assigned exactly **one label** from a predefined set of classes, but it is still considered a "single-label" classification problem. The main difference between multiclass and singlelabel is semantic. **Multiclass** is more commonly used when you have multiple classes, but you still need just **one class label per image**.

○ **Multilabel**

Explanation

Multilabel classification involves assigning **multiple labels** to a single image. This is used when an image can belong to more than one category (for example, if an image can have both a dog and a cat).

○ **Singleclass**

Explanation

Singleclass is not a valid classification term used in Azure AI Vision. It might be a confusion between **singlelabel** or a misrepresentation of a concept. The term **singleclass** would imply that the model can only classify images into one single class, but that doesn't properly describe how classification works in Azure AI Vision.

Correct answer

◯ **Singlelabel**

Explanation

The term **"singlelabel"** is used when an image is assigned to a single class from a set of possible classes. This approach ensures that each image is classified with **only one label**, which fits the scenario where you need to classify each image with just one animal. This ensures that there is only one class (or label) per image, which matches the requirement of having a single animal per image.

◯ **Unilabel**

Explanation

Unilabel is not a recognized term in Azure AI Vision classification terminology. It is likely an incorrect or confused version of **singlelabel**. The right term for classifying one label per image from a set of classes is **singlelabel**, not unilabel.

Question 16:

You are developing an application named "VisionApp" that uses Azure's Image Analysis API to extract meaningful data from images.

You send a request with the following endpoint:

```
https://<resource-name>.cognitiveservices.azure.com/computervision/imageanalysis:analyze?features=read,description
```

What will be the result of this API request?

- ○ **The objects detected in the image and their approximate locations**
- ○ **A description of the image content only**
- ◉ **The visible text in the image and a description of the image content**

You are developing a customer support chatbot using Azure OpenAI Service. To optimize the chatbot's responses during testing, you decide to follow Microsoft's best practices for prompt engineering.

Which three strategies should you apply to create effective prompts for accurate and relevant answers?

☐ **Use Minimal Text**

Explanation

1. While keeping prompts concise is useful, minimal text can omit critical context or details.
2. **Why it's incorrect:** Omitting essential information leads to incomplete answers.

☐ **Be Vague**

Explanation

1. Vague prompts lack clarity and result in incomplete or irrelevant responses.
2. **Why it's incorrect:** AI models require clear, specific instructions to generate high-quality answers.

Your selection is correct

☑ **Provide Context**

Explanation

1. Including relevant context in the prompt ensures the AI understands the question's background.
2. **Why it's correct:** Context helps the model generate more targeted and meaningful responses.

Your selection is correct

☑ **Be Clear and Specific**

Explanation

1. Using clear and specific prompts avoids ambiguity and ensures precise answers.

2. **Why it's correct:** Specific prompts reduce the chance of irrelevant responses.

Your selection is correct

☑ **Order Instructions Logically**

Explanation

1. Structuring instructions logically helps the AI interpret the query correctly.
2. **Why it's correct:** The sequence of instructions affects how the model processes and generates answers.

Question 18:

You are designing a knowledge-mining solution using **Azure Cognitive Search**. During the enrichment pipeline process, you need to extract text content from a file such as PDFs, Word documents, or images.

Which built-in cognitive skill should you configure to extract this content?

- (●) **Microsoft.Skills.Util.DocumentExtractionSkill**
- ( ) **Microsoft.Skills.Text.V3.EntityRecognitionSkill**
- ( ) **Microsoft.Skills.Text.SplitSkill**
- ( ) **Microsoft.Skills.Text.KeyPhraseExtractionSkill**

Question 19:

You are developing a web app that will use the DALL-E 3 model via Azure OpenAI to generate images based on user prompts. The app needs to send HTTP requests to the Azure OpenAI API to ensure the image generation process works correctly.

Which three HTTP header properties should be included in the API requests to ensure successful image generation?

- [ ] **User's prompt**
- [x] **Azure OpenAI service resource name**
- [x] **API version used in the operation**
- [ ] **Image resolution settings**
- [ ] **Desired image style**
- [x] **DALL-E 3 model deployment name**

You are developing a video processing application that will utilize Azure AI Video Indexer to extract insights from videos containing content in multiple languages.

To configure the API calls for identifying and processing the multilingual content, you need to specify the correct parameter for language identification.

Which value should you set for the `sourceLanguage` parameter to ensure accurate multilingual detection?

- ○ **multi-lingual detection**

Explanation

Although similar in meaning to "multi-language detection," "multi-lingual detection" is not the correct API parameter value. The Azure AI Video Indexer uses the term **"multi-language detection"** to refer to the feature that supports identifying and processing multiple languages. The use of "multi-lingual" would not be recognized as a valid configuration value in this context.

Correct answer

- ○ **multi-language detection**

Explanation

The "multi-language detection" value is specifically used in the Azure AI Video Indexer API to enable the model to detect multiple languages within the video content. When working with videos containing more than one language, this parameter ensures that the system can identify and process each language properly, enabling it to extract insights across multiple language tracks.

- ○ **language detection**

Explanation

The term "language detection" is too generic and does not explicitly indicate multilingual support. In Azure AI Video Indexer, the value "language detection" typically refers to detecting a single language in the content. This would not provide the needed functionality for multi-language content.

Your answer is incorrect

- ⊗ **multi-detection**

Question 21:

You are developing a web app named **App1** hosted in **Azure App Services**. The app must securely authenticate to **Azure AI Services** using **Microsoft Entra ID** while following these principles:

- Minimize administrative effort
- Apply the principle of least privilege

What should you do?

○ **reate a Microsoft Entra app registration and enable certificate-based authentication**

Explanation

1. While **app registration** with **certificate-based authentication** can secure App1, managing certificates adds **administrative overhead.**
2. Certificates must be created, renewed, and secured properly, making this option **less efficient** compared to managed identities.

Correct answer

○ **From App1, enable a managed identity and assign role-based access control (RBAC) permissions to Azure AI Services**

Explanation

1. **Managed identities** for Azure resources enable secure, password-less authentication to Azure services, including **Azure AI Services.**
2. It **minimizes administrative effort** by removing the need to manage secrets or certificates.
3. By assigning specific **RBAC roles**, access can be limited to only the **required permissions**, ensuring **least privilege** access.

Your answer is incorrect

⊗ **Create a secret and store the secret in an Azure Key Vault. Assign App1 role-based access control (RBAC) permissions to the secre**

Explanation

1. Storing a secret in **Azure Key Vault** can secure credentials but **increases**

Question 22:

You are designing a solution using **Azure Cognitive Search** and need to define a **skillset** to perform data enrichment tasks.

- ○ **name, description, knowledgeStore, and encryptionKey**

Explanation

1. `knowledgeStore:` This is an optional feature that persists enriched data in Azure storage. It is not mandatory for defining a skillset unless data persistence outside the index is required.
2. `encryptionKey:` This is optional and used for securing sensitive data during indexing. Including it is useful but not required by default.

Correct answer

- ○ **name, description, and skills**

Explanation

1. `name:`
   - A unique identifier for the skillset within Azure Cognitive Search. It distinguishes this skillset from others in the same service.
2. `description:`
   - A brief explanation of the skillset's purpose. Though technically optional in some contexts, it is recommended for clarity.
3. `skills:`
   - A list of skill definitions that specify how to extract, transform, and enrich data. Each skill defines input and output fields, as well as processing logic.

- ○ **name, description, skills, knowledgeStore, and encryptionKey**

Explanation

This option includes unnecessary fields such as `knowledgeStore` and `encryptionKey`, which are both optional

Your answer is incorrect

⊗ **name, description, skills, and cognitiveServices**

Explanation

- `cognitiveServices:` This defines whether Azure Cognitive Services should be used with the skillset. However, it is not required if you are using skills that don't rely on cognitive services or if you use custom skills.

Question 23:

You are designing an Azure AI-powered assistant using a generative AI model.

You want to guide the model's responses by configuring its system message during prompt initialization.

Which two functions can a system message perform when configuring the model?

- [x] **Sets the rules for how the assistant should behave**
- [ ] **Determines the maximum size of data the assistant can process**
- [ ] **Automatically filters out sensitive data from input prompts**
- [x] **Defines the assistant's tone and interaction style**

Question 24:

You are developing an AI-powered mobile app for an edge device, which will analyze images of fruits using Azure AI Custom Vision. The app needs to run without internet connectivity and should be deployed on the edge device.

Which model domain should you select to ensure that the app works offline?

Correct answer

○ **Compact domain**

Explanation

Compact domain is the correct choice because it is specifically designed to support running AI models on edge devices without requiring an internet connection. The "compact" domain is optimized for offline use, meaning that the model can be deployed directly on the edge device for real-time image analysis, even when the device is not connected to the internet.

Your answer is incorrect

⊗ **Food domain**

Explanation

The **Food domain** is a specialized domain in Azure AI Custom Vision designed to classify food items, but it is not specifically designed for offline use. It's primarily intended for cloud-based deployments where internet access is available, and it may not be suitable for running completely offline, especially in edge scenarios.

- **Why Incorrect:** While the food domain is specialized for identifying food items, it does not provide the offline functionality that the "compact domain" offers, making it unsuitable for edge devices that need to run without internet access.

○ **General [A2] domain**

Explanation

Like the General [A1] domain, the General [A2] domain is not suitable for offline operations on edge devices. It relies on cloud infrastructure, which makes it unsuitable for the specific requirement of offline functionality.

○ **General [A1] domain**

Explanation

The **General [A1] domain** is a general-purpose model, but it is intended for cloud-based image classification. Like the food domain, it requires an internet connection to process and analyze images.

- **Why Incorrect:** This domain is not optimized for offline usage on edge devices, so it wouldn't be able to run without internet connectivity.

You are developing a knowledge mining solution using Azure AI Search, and you need to support advanced search capabilities, including the use of wildcard queries.

Which option should you include in the REST API request to enable wildcard queries in search requests?

- ○ **"queryType": "advanced"**

Explanation

There is no specific `"advanced"` query type in Azure AI Search's REST API. Azure AI Search only supports `"simple"`, `"full"`, and `"extended"` query types. The `"advanced"` term might sound logical, but it is not recognized in the context of Azure AI Search.

- ○ **"queryType": "simple"**

Explanation

The `"simple"` query type enables basic, simple search functionality without supporting advanced features like wildcard queries. It is a lightweight option for searches that don't need complex querying capabilities, so it does not allow wildcard searches.

Correct answer

○ **"queryType": "extended"**

Explanation

The `"extended"` query type allows for advanced query features in Azure AI Search, including wildcard searches. Wildcard queries let you match patterns in the text fields (e.g., using `*` to replace multiple characters or `?` to replace a single character). To enable wildcard queries, you must explicitly specify `"queryType": "extended"` in the REST API request. This setting activates the broader querying capabilities that support wildcard searches along with other advanced search features like fuzzy matching, proximity searches, and more.

Your answer is incorrect

⊗ **"queryType": "full"**

Explanation

The `"full"` query type is used to enable the full-text search capabilities of Azure AI Search. However, it does not specifically support wildcard queries. The `"full"` query type is more general and allows for complex queries, but to specifically use wildcard characters in queries, you need the `"extended"` query type.